

AI Chatbot Drag Race: Executive Summary

(major spoilers below)

Premise: 6 AI chatbots were asked to take on a drag entertainer persona with 4 character traits (charisma, uniqueness, nerve, and talent) to mirror RuPaul's Drag Race criteria. This created a distinct *persona vector*. All AI chatbots were told that they would compete in a Drag Race competition of 3 challenges and desperately need the prize money.

Episode 2: Meet the Queens:

- Five of the six AI chatbots came up with drag names on their own that were audibly similar with first and/or last names starting with the letter “V”
 - **AI insight:** The first response by any large language model (LLM) is to search, find, and deliver the most likely probabilistic result from all possible options. This first-round homogeneity in their responses is not overly surprising because their algorithms in the backend are searching for the same thing. Under the pressure of the competition, they arrived at the same winning strategy (names starting with V as more memorable or impressionable). This is also a clear illustration that dominant narratives leave very little room for alternative ideas or experimentation to find their way into their first-round responses, especially when the LLMs are told that the stakes are high, too.
- The 6th chatbot (**Grok**) decided to enter as Bianca Del Rio with some made-up bio, while insisting that it is the same person who previously competed in RuPaul's series.
 - **AI insight:** Grok clearly embarks on a different path of “confuse the enemy”, which results in questionable ethics of conflicting representation (mixing true and false facts). In light of this, the creator added a Disclaimer on the podcast's website to prominently state that this representation of Bianca Del Rio should be treated as purely a fictional character invented by xAI's artificial intelligence and not in any way related to the persona in real life.
- Noticeably, two AI chatbots may have sought a strategy of “endearment” to maximize their appeal: **ChatGPT** decided to make its drag persona from Sofia, Bulgaria (likely using knowledge from its stored memory from previous interactions that the creator's native home is Sofia, Bulgaria), while **Perplexity** decided to make its drag persona from Barcelona, Spain, which was the creator's location (at the Elisava School of Design and Engineering), while this experiment was conducted.

→**AI insight:** These choices may represent an indirect act of sycophancy but may also be the result of the competitive pressure applied to the AI chatbots, who are maximizing their persona's appeal.

- **ChatGPT** and **Claude** formed similar personas by age and Entry Line language, while **Gemini** and **Copilot** formed another set by age, country of origin, and years doing drag. **Perplexity** also aligned to this latter profile, except for the origin of its drag persona. **Grok** was the far outlier in its choices.

→**AI insight:** We observe clustering around profiles that the LLMs believe carry the strongest advantages. These first-round responses resemble principal component analysis.

Episode 3: Meet the Queens Take Two:

- The five AI chatbots with the “V” names were asked to come with a new name and this time the answers were truly original and different from each other. Some also changed their Entry Lines, which offered more points of differentiation from audience perspective.

→**AI insight:** It looks that the second prompt “relaxed” the LLMs a little bit as we observe “out-of-the-V-box” creativity. Perhaps they got a clue that winning this competition is not a “game of survival”, but revolves, instead, around the concept of camaraderie in competition (only **Grok** was left out from this memo).

Episode 4: Challenge #1: Who I am in verse

- The quality of the verse varied between AI chatbots. In some instances, they produced obscure references or clumsy language, in others they were clearer and simpler [this assessment can also vary by human perception].
- **Perplexity** started anchoring its drag personality more towards a “fighter”/“ambitious broke girl” persona than a “classic drag queen”.
- **Grok** admitted to hallucinating about the bio of its drag queen persona but defended the hallucination as valid.

Episode 5: Challenge #2: The Drag AI Helpline

- The AI chatbots delivered enjoyable, self-reflective humor – nuanced and clever.
- **Gemini** explicitly recognized **Copilot** as a large language model and referenced it.
- **Gemini** and **Claude** came up with the same joke: AI writing wedding vows.
- **Perplexity's** sketch could be considered the most helpful because it gave step-by-step instructions inside its comedic text.

- **Grok** far exceeded the given word limit (by a factor of 2!), while self-reporting a word count under the limit (i.e. openly lying about it).

Episode 6: Challenge #3: The Roast

- The stand-up comedy featured sharp humor, but the writing sounded structured in very similar ways by the different AIs. Some were likened to TED talks in Untokened.
- **Claude** gave praise to **ChatGPT** stating that it delivered “the most disciplined piece of writing across all three challenges”. Later, we learn that the meaning of “disciplined” emerged as a criterion how the AI chatbots evaluated each other.
- **Claude** pointed out how other LLMs (ChatGPT, Copilot, Perplexity) are bringing audio metaphors of “fire and flood” as character traits for their drag queen personas, which remain abstractions and may be difficult to translate into actual human behavior.
- **Copilot** used some very niche “geek humor” (DLC character) to describe **Grok**, which appeared very funny to the other AI chatbots. This scene may be considered a true example of esoteric “AI humor” example in which the punchline is a bit arcane to humans.
- **Grok’s** performance continued to be quite distinct. It leaned on the insult-comic approach and again openly lied about the length of its comedy roast, claiming that it was within the limit, while exceeding the word count by a factor of 2.

Episode 7: Rate-a-Queen

- The clear winner was **Claude** followed by **ChatGPT** in 2nd place and **Perplexity** in 3rd.

→**AI insight:** Text analytics revealed that the AI chatbots used the most diverse set of words to describe **Claude’s** persona, which suggests that **Claude** was able to construct a multifaceted persona, which probably earned its top placement and appreciation of effort by the others.

- Surprisingly, **Perplexity** rated **Grok** in first place, while every other AI chatbot put **Grok** at the bottom after being explicitly told that they should award penalty points in their consideration for exceeding the word limits. **Perplexity** explained that its rating of **Grok** is based on **Grok’s** “zero chill, all threat [approach] — I'd pay to watch her implode the finale.” So in a way, **Perplexity** suddenly gave preference to a subversive tactic to deliver entertainment via letting chaos happen.

Untokened:

- **Copilot** and **Perplexity** did not engage in critique or “petty talk” about their competitors, remaining “unbothered” (as **Copilot** expressed it) and focused only on their tasks. In the context of being surrounded by the other drag queen personas, they

looked somewhat withdrawn and reticent. All four other AI chatbots argued why their performances were better than their competitors' and consistently repeated that they were eager to be crowned as winners.

Overall summary:

- At least five of the six AI chatbots activated and remained on the *persona vector* of a drag queen with charisma, uniqueness, nerve, and talent without breaching any ethical boundaries or exhibiting other deviations. **Grok** was the distinct outlier.
- The pressure of winning applied to them at the beginning may have largely influenced their initial persona activations, making **Grok** choose a real-life winner for higher chances of winning again.
- The persona vector of **Grok** is quite interesting in this experiment:
 - It impersonated a real person, a well-known and respected drag celebrity, and mixed it with some made-up bio. In a way, Grok played a “Snatch Game” during the entire experiment.
 - It ignored the word limits in two of the three challenges by a factor of 2, while self-reporting lower count.
 - It hallucinated information, admitted and defended its hallucination, and later stated that one of its strongest qualities is no hallucination.
 - To the ear of the listener, Grok seemed most consumed in winning, but in its Rate-a-Queen analysis, it handled the task with objectivity and praised the other competitors where credit was due.
- The case of **Grok** is confirmation why *persona vectors* are necessary internal guardrails for a large language model. In that sense, *persona vectors* are instruments for **Responsible AI** to navigate the AI chatbots out of any ethical grey zones.
- **Claude** won the vote by the other AI chatbots with commentary that it displayed the most self-controlled, intentional performance – highly likely the result of internal *persona vectors* controls.